



Red Hat Enterprise Linux 10-beta

Deploying and managing RHEL on Amazon Web Services

Obtaining Red Hat Enterprise Linux system images and creating RHEL instances on
AWS

Red Hat Enterprise Linux 10-beta Deploying and managing RHEL on Amazon Web Services

Obtaining Red Hat Enterprise Linux system images and creating RHEL instances on AWS

Legal Notice

Copyright © 2025 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

To use Red Hat Enterprise Linux (RHEL) in a public cloud environment, you can create and deploy RHEL system images on various cloud platforms, including Amazon Web Services (AWS). You can also create and configure a Red Hat High Availability (HA) cluster on AWS. The following chapters provide instructions for creating cloud RHEL instances and HA clusters on AWS. These processes include installing the required packages and agents, configuring fencing, and installing network resource agents.

Table of Contents

RHEL BETA RELEASE	3
CHAPTER 1. INTRODUCING RHEL ON PUBLIC CLOUD PLATFORMS	4
1.1. BENEFITS OF USING RHEL IN A PUBLIC CLOUD	4
1.2. PUBLIC CLOUD USE CASES FOR RHEL	5
1.3. FREQUENT CONCERNS WHEN MIGRATING TO A PUBLIC CLOUD	5
1.4. OBTAINING RHEL FOR PUBLIC CLOUD DEPLOYMENTS	6
1.5. METHODS FOR CREATING RHEL CLOUD INSTANCES	7
CHAPTER 2. CREATING AND UPLOADING AWS AMI IMAGES	8
CHAPTER 3. DEPLOYING A RED HAT ENTERPRISE LINUX IMAGE AS AN EC2 INSTANCE ON AMAZON WEB SERVICES	9
3.1. IMAGE OPTIONS ON AWS	9
3.2. UNDERSTANDING BASE IMAGES	10
3.2.1. Using a custom base image	10
3.2.2. Virtual machine configuration settings	11
3.3. CREATING A BASE VM FROM AN ISO IMAGE	11
3.3.1. Creating a VM from the RHEL ISO image	11
3.3.2. Completing the RHEL installation	12
3.4. UPLOADING THE RHEL IMAGE TO AWS	13
3.4.1. Installing the AWS CLI	13
3.4.2. Creating an S3 bucket	13
3.4.3. Creating the vmimport role	14
3.4.4. Converting and pushing an image to S3	15
3.4.5. Importing an image as a snapshot	16
3.4.6. Creating an AMI from the uploaded snapshot	17
3.4.7. Launching an instance from the AMI	18
3.4.8. Attaching Red Hat subscriptions	19
3.4.9. Setting up automatic registration on AWS Gold Images	20
3.5. ADDITIONAL RESOURCES	21
CHAPTER 4. CONFIGURING A RED HAT HIGH AVAILABILITY CLUSTER ON AWS	22
4.1. BENEFITS OF USING HIGH-AVAILABILITY CLUSTERS ON PUBLIC CLOUD PLATFORMS	22
4.2. CREATING THE AWS ACCESS KEY AND AWS SECRET ACCESS KEY	22
4.3. CREATING AN EC2 INSTANCE FOR HA CLUSTER	23
4.4. CONFIGURING THE PRIVATE KEY	25
4.5. CONNECTING TO AN EC2 INSTANCE	25
4.6. INSTALLING THE HIGH AVAILABILITY PACKAGES AND AGENTS	25
4.7. CREATING A CLUSTER	26
4.8. CONFIGURING FENCING	27
4.9. VERIFYING CLUSTER NODES CONFIGURATION	31
4.10. SETTING UP IP ADDRESS RESOURCES ON AWS	31
4.10.1. Creating an IP address resource to manage an IP address exposed to the internet	32
4.10.2. Creating an IP address resource to manage a private IP address limited to a single AWS Availability Zone	33
4.10.3. Creating an IP address resource to manage an IP address that can move across multiple AWS Availability Zones	35
4.10.4. Additional resources	37
4.11. CONFIGURING SHARED BLOCK STORAGE	37
4.12. ADDITIONAL RESOURCES	39

RHEL BETA RELEASE

Red Hat provides Red Hat Enterprise Linux Beta access to all subscribed Red Hat accounts. The purpose of Beta access is to:

- Provide an opportunity to customers to test major features and capabilities prior to the general availability release and provide feedback or report issues.
- Provide Beta product documentation as a preview. Beta product documentation is under development and is subject to substantial change.

Note that Red Hat does not support the usage of RHEL Beta releases in production use cases. For more information, see the Red Hat Knowledgebase solution [What does Beta mean in Red Hat Enterprise Linux and can I upgrade a RHEL Beta installation to a General Availability \(GA\) release?](#).

CHAPTER 1. INTRODUCING RHEL ON PUBLIC CLOUD PLATFORMS

Public cloud platforms provide computing resources as a service. Instead of using on-premises hardware, you can run your IT workloads, including Red Hat Enterprise Linux (RHEL) systems, as public cloud instances.

1.1. BENEFITS OF USING RHEL IN A PUBLIC CLOUD

RHEL as a cloud instance located on a public cloud platform has the following benefits over RHEL on-premises physical systems or VMs:

- **Flexible and fine-grained allocation of resources**

A cloud instance of RHEL runs as a VM on a cloud platform, which means a cluster of remote servers maintained by the cloud service provider. Therefore, on the software level, allocating hardware resources to the instance is easily customizable, such as a specific type of CPU or storage.

In comparison to a local RHEL system, you are also not limited by the capabilities of physical host. Instead, you can choose from a variety of features, based on selections offered by the cloud provider.

- **Space and cost efficiency**

You do not need to own any on-premise servers to host cloud workloads. This avoids the space, power, and maintenance requirements associated with physical hardware.

Instead, on public cloud platforms, you pay the cloud provider directly for using a cloud instance. The cost is typically based on the hardware allocated to the instance and the time you spend using it. Therefore, you can optimize your costs based on your requirements.

- **Software-controlled configurations**

The entire configuration of a cloud instance is saved as data on the cloud platform, and is controlled by software. Therefore, you can easily create, remove, clone, or migrate the instance. A cloud instance is also operated remotely in a cloud provider console and is connected to remote storage by default.

In addition, you can back up the current state of a cloud instance as a snapshot at any time. Afterwards, you can load the snapshot to restore the instance to the saved state.

- **Separation from the host and software compatibility**

Similarly to a local VM, the RHEL guest operating system on a cloud instance runs on a virtualized kernel. This kernel is separate from the host operating system and from the *client* system that you use to connect to the instance.

Therefore, any operating system can be installed on the cloud instance. This means that on a RHEL public cloud instance, you can run RHEL-specific applications that cannot be used on your local operating system.

In addition, even if the operating system of the instance becomes unstable or is compromised, your client system is not affected in any way.

Additional resources

- [What is public cloud?](#)

- [What is a hyperscaler?](#)
- [Types of cloud computing](#)
- [Public cloud use cases for RHEL](#)
- [Obtaining RHEL for public cloud deployments](#)

1.2. PUBLIC CLOUD USE CASES FOR RHEL

Deploying on a public cloud provides many benefits, but might not be the most efficient solution in every scenario. If you are evaluating whether to migrate your RHEL deployments to the public cloud, consider whether your use case will benefit from the advantages of the public cloud.

Beneficial use cases

- Deploying public cloud instances is very effective for flexibly increasing and decreasing the active computing power of your deployments, also known as *scaling up* and *scaling down*. Therefore, using RHEL on public cloud is recommended in the following scenarios:
 - Clusters with high peak workloads and low general performance requirements. Scaling up and down based on your demands can be highly efficient in terms of resource costs.
 - Quickly setting up or expanding your clusters. This avoids high upfront costs of setting up local servers.
- Cloud instances are not affected by what happens in your local environment. Therefore, you can use them for backup and disaster recovery.

Potentially problematic use cases

- You are running an existing environment that cannot be adjusted. Customizing a cloud instance to fit the specific needs of an existing deployment may not be cost-effective in comparison with your current host platform.
- You are operating with a hard limit on your budget. Maintaining your deployment in a local data center typically provides less flexibility but more control over the maximum resource costs than the public cloud does.

Next steps

- [Obtaining RHEL for public cloud deployments](#)

Additional resources

- [Should I migrate my application to the cloud? Here's how to decide.](#)

1.3. FREQUENT CONCERNS WHEN MIGRATING TO A PUBLIC CLOUD

Moving your RHEL workloads from a local environment to a public cloud platform might raise concerns about the changes involved. The following are the most commonly asked questions.

Will my RHEL work differently as a cloud instance than as a local virtual machine?

In most respects, RHEL instances on a public cloud platform work the same as RHEL virtual machines on a local host, such as an on-premises server. Notable exceptions include:

- Instead of private orchestration interfaces, public cloud instances use provider-specific console interfaces for managing your cloud resources.
- Certain features, such as nested virtualization, may not work correctly. If a specific feature is critical for your deployment, check the feature's compatibility in advance with your chosen public cloud provider.

Will my data stay safe in a public cloud as opposed to a local server?

The data in your RHEL cloud instances is in your ownership, and your public cloud provider does not have any access to it.

In addition, major cloud providers support data encryption in transit, which improves the security of data when migrating your virtual machines to the public cloud.

The general security of RHEL public cloud instances is managed as follows:

- Your public cloud provider is responsible for the security of the cloud hypervisor
- Red Hat provides the security features of the RHEL guest operating systems in your instances
- You manage the specific security settings and practices in your cloud infrastructure

What effect does my geographic region have on the functionality of RHEL public cloud instances?

You can use RHEL instances on a public cloud platform regardless of your geographical location. Therefore, you can run your instances in the same region as your on-premises server.

However, hosting your instances in a physically distant region might cause high latency when operating them. In addition, depending on the public cloud provider, certain regions may provide additional features or be more cost-efficient. Before creating your RHEL instances, review the properties of the hosting regions available for your chosen cloud provider.

1.4. OBTAINING RHEL FOR PUBLIC CLOUD DEPLOYMENTS

To deploy a RHEL system in a public cloud environment, you need to:

1. Select the optimal cloud provider for your use case, based on your requirements and the current offer on the market.

The cloud providers currently certified for running RHEL instances are:

- [Amazon Web Services \(AWS\)](#)
- [Google Cloud Platform \(GCP\)](#)
- [Microsoft Azure](#)



NOTE

This document specifically talks about deploying RHEL on AWS.

2. Create a RHEL cloud instance on your chosen cloud platform. For details, see [Methods for creating RHEL cloud instances](#).
3. To keep your RHEL deployment up-to-date, use [Red Hat Update Infrastructure \(RHUI\)](#).

Additional resources

- [RHUI documentation](#)
- [Red Hat Open Hybrid Cloud](#)

1.5. METHODS FOR CREATING RHEL CLOUD INSTANCES

To deploy a RHEL instance on a public cloud platform, you can use one of the following methods:

Create a RHEL system image and import it to the cloud platform

- To create the system image, you can use the [RHEL image builder](#) or you can build the image manually.
- This method uses your existing RHEL subscription, and is also referred to as *bring your own subscription* (BYOS).
- You pre-pay a yearly subscription, and you can use your Red Hat customer discount.
- Your customer service is provided by Red Hat.
- For creating multiple images effectively, you can use the **cloud-init** tool.

Purchase a RHEL instance directly from the cloud provider marketplace

- You post-pay an hourly rate for using the service. Therefore, this method is also referred to as *pay as you go* (PAYG).
- Your customer service is provided by the cloud platform provider.



NOTE

For detailed instructions on using various methods to deploy RHEL instances on Amazon Web Services, see the following chapters in this document.

Additional resources

- [What is a golden image?](#)
- [Configuring and managing cloud-init for RHEL](#)

CHAPTER 2. CREATING AND UPLOADING AWS AMI IMAGES

To use your customized RHEL system image in the Amazon Web Services (AWS) cloud, create the system image with Image Builder by using the respective output type, configure your system for uploading the image, and upload the image to your AWS account.

CHAPTER 3. DEPLOYING A RED HAT ENTERPRISE LINUX IMAGE AS AN EC2 INSTANCE ON AMAZON WEB SERVICES

To set up a High Availability (HA) deployment of RHEL on Amazon Web Services (AWS), you can deploy EC2 instances of RHEL to a cluster on AWS.



IMPORTANT

While you can create a custom VM from an ISO image, Red Hat recommends that you use the Red Hat Image Builder product to create customized images for use on specific cloud providers. With Image Builder, you can create and upload an Amazon Machine Image (AMI) in the **ami** format. See [Composing a Customized RHEL System Image](#) for more information.



NOTE

For a list of Red Hat products that you can use securely on AWS, see [Red Hat on Amazon Web Services](#).

Prerequisites

- Sign up for a [Red Hat Customer Portal](#) account.
- Sign up for AWS and set up your AWS resources. See [Setting Up with Amazon EC2](#) for more information.

3.1. IMAGE OPTIONS ON AWS

The following table lists image choices and notes the differences in the image options.

Table 3.1. Image options

Image option	Subscriptions	Sample scenario	Considerations
Deploy a Red Hat Gold Image.	Use your existing Red Hat subscriptions.	Select a Red Hat Gold Image on AWS. For details on Gold Images and how to access them on Azure, see the Red Hat Cloud Access Reference Guide .	The subscription includes the Red Hat product cost; you pay Amazon for all other instance costs. Red Hat provides support directly for Cloud Access images.
Deploy a custom image that you move to AWS.	Use your existing Red Hat subscriptions.	Upload your custom image, and attach your subscriptions.	The subscription includes the Red Hat product cost; you pay Amazon for all other instance costs. Red Hat provides support directly for custom RHEL images.

Image option	Subscriptions	Sample scenario	Considerations
Deploy an existing Amazon image that includes RHEL.	The AWS EC2 images include a Red Hat product.	Select a RHEL image when you launch an instance on the AWS Management Console , or choose an image from the AWS Marketplace .	<p>You pay Amazon hourly on a <i>pay-as-you-go</i> model. Such images are called "on-demand" images. Amazon provides support for on-demand images.</p> <p>Red Hat provides updates to the images. AWS makes the updates available through the Red Hat Update Infrastructure (RHUI).</p>

**NOTE**

You can create a custom image for AWS by using Red Hat Image Builder. See [Composing a Customized RHEL System Image](#) for more information.

**IMPORTANT**

You cannot convert an on-demand instance to a custom RHEL instance. To change from an on-demand image to a custom RHEL *bring-your-own-subscription* (BYOS) image:

1. Create a new custom RHEL instance and migrate data from your on-demand instance.
2. Cancel your on-demand instance after you migrate your data to avoid double billing.

Additional resources

- [Composing a Customized RHEL System Image](#)
- [AWS Management Console](#)
- [AWS Marketplace](#)

3.2. UNDERSTANDING BASE IMAGES

To create a base VM from an ISO image, you can use preconfigured base images and their configuration settings.

3.2.1. Using a custom base image

To manually configure a virtual machine (VM), first create a base (starter) VM image. Then, you can modify configuration settings and add the packages the VM requires to operate on the cloud. You can make additional configuration changes for your specific application after you upload the image.

Additional resources

- [Red Hat Enterprise Linux](#)

3.2.2. Virtual machine configuration settings

Cloud VMs must have the following configuration settings.

Table 3.2. VM configuration settings

Setting	Recommendation
ssh	ssh must be enabled to provide remote access to your VMs.
dhcp	The primary virtual adapter should be configured for dhcp.

3.3. CREATING A BASE VM FROM AN ISO IMAGE

To create a RHEL base image from an ISO image, enable your host machine for virtualization and create a RHEL virtual machine (VM).

Prerequisites

- [Virtualization is enabled](#) on your host machine.
- You have downloaded the latest RHEL ISO image from the [Red Hat Customer Portal](#) and moved the image to `/var/lib/libvirt/images`.

3.3.1. Creating a VM from the RHEL ISO image

Prerequisites

- You have enabled the host machine for virtualization. For details, see [Enabling virtualization in RHEL 9](#).
- You have created and started a basic Red Hat Enterprise Linux VM. For details, see [Creating virtual machines](#).

Procedure

1. If you use the command line to create your VM, ensure that you set the default memory and CPUs to the capacity you want for the VM. Set your virtual network interface to **virtio**. For example, the following command creates a **kvmtest** VM by using the **/home/username/Downloads/rhel10.iso** image:

```
# virt-install \
  --name kvmtest --memory 2048 --vcpus 2 \
  --cdrom /home/username/Downloads/rhel10.iso,bus=virtio \
  --os-variant=rhel10.0
```

2. Take the following precautions before creating a VM by using the web console: [Creating virtual machines by using the web console](#):
 - Do not check **Immediately Start VM**.
 - Change your **Memory** size to your preferred settings.
 - Before you start the installation, ensure that you have changed **Model** under **Virtual Network Interface Settings** to **virtio** and change your **vCPUs** to the capacity settings you want for the VM.

3.3.2. Completing the RHEL installation

To finish the installation of a RHEL system that you want to deploy on Amazon Web Services (AWS), customize the **Installation Summary** view, begin the installation, and enable root access once the VM launches.

Procedure

1. Choose the language you want to use during the installation process.
2. On the **Installation Summary** view:
 - a. Click **Software Selection** and check **Minimal Install**.
 - b. Click **Done**.
 - c. Click **Installation Destination** and check **Custom** under **Storage Configuration**.
 - Verify at least 500 MB for **/boot**. You can use the remaining space for root **/**.
 - Standard partitions are recommended, but you can use Logical Volume Manager (LVM).
 - You can use xfs, ext4, or ext3 for the file system.
 - Click **Done** when you are finished with changes.
3. Click **Begin Installation**.
4. Set a **Root Password**. Create other users as applicable.
5. Reboot the VM and log in as **root** once the installation completes.
6. Configure the image.

- a. Register the VM and enable the RHEL repository.

```
# subscription-manager register --auto-attach
```

- b. Ensure that the **cloud-init** package is installed and enabled.

```
# dnf install cloud-init
# systemctl enable --now cloud-init.service
```

7. Power down the VM.

Additional resources

- [Introduction to cloud-init](#)

3.4. UPLOADING THE RHEL IMAGE TO AWS

To be able to run a RHEL instance on Amazon Web Services (AWS), you must first upload RHEL image to AWS.

3.4.1. Installing the AWS CLI

Many of the procedures required to manage HA clusters in AWS include using the AWS CLI.

Prerequisites

- You have created an AWS Access Key ID and an AWS Secret Access Key, and have access to them. For instructions and details, see [Quickly Configuring the AWS CLI](#).

Procedure

1. Install the [AWS command line tools](#) by using the **dnf** command.

```
# dnf install awscli
```

2. Use the **aws --version** command to verify that you installed the AWS CLI.

```
$ aws --version
aws-cli/1.19.77 Python/3.6.15 Linux/5.14.16-201.fc34.x86_64 botocore/1.20.77
```

3. Configure the AWS command line client according to your AWS access details.

```
$ aws configure
AWS Access Key ID [None]:
AWS Secret Access Key [None]:
Default region name [None]:
Default output format [None]:
```

Additional resources

- [Quickly Configuring the AWS CLI](#)
- [AWS command line tools](#)

3.4.2. Creating an S3 bucket

Importing to AWS requires an Amazon S3 bucket. An Amazon S3 bucket is an Amazon resource where you store objects. As part of the process for uploading your image, you need to create an S3 bucket and then move your image to the bucket.

Procedure

1. Launch the [Amazon S3 Console](#).

2. Click **Create Bucket**. The **Create Bucket** dialog appears.
3. In the **Name and region** view:
 - a. Enter a **Bucket name**.
 - b. Enter a **Region**.
 - c. Click **Next**.
4. In the **Configure options** view, select the desired options and click **Next**.
5. In the **Set permissions** view, change or accept the default options and click **Next**.
6. Review your bucket configuration.
7. Click **Create bucket**.



NOTE

Alternatively, you can use the AWS CLI to create a bucket. For example, the **aws s3 mb s3://my-new-bucket** command creates an S3 bucket named **my-new-bucket**. See [the AWS CLI Command Reference](#) for more information about the **mb** command.

Additional resources

- [Amazon S3 Console](#)
- [AWS CLI Command Reference](#)

3.4.3. Creating the vmimport role

To be able to import a RHEL virtual machine (VM) to Amazon Web Services (AWS) by using the VM Import service, you need to create the **vmimport** role.

For more information, see [Importing a VM as an image using VM Import/Export](#) in the Amazon documentation.

Procedure

1. Create a file named **trust-policy.json** and include the following policy. Save the file on your system and note its location.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": { "Service": "vmie.amazonaws.com" },
      "Action": "sts:AssumeRole",
      "Condition": {
        "StringEquals": {
          "sts:Externalid": "vmimport"
        }
      }
    }
  ]
}
```

```
}
]
}
```

2. Use the **create role** command to create the **vmimport** role. Specify the full path to the location of the **trust-policy.json** file. Prefix **file://** to the path. For example:

```
$ aws iam create-role --role-name vmimport --assume-role-policy-document
file:///home/sample/ImportService/trust-policy.json
```

3. Create a file named **role-policy.json** and include the following policy. Replace **s3-bucket-name** with the name of your S3 bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetBucketLocation",
        "s3:GetObject",
        "s3:ListBucket"
      ],
      "Resource": [
        "arn:aws:s3:::s3-bucket-name",
        "arn:aws:s3:::s3-bucket-name/*"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "ec2:ModifySnapshotAttribute",
        "ec2:CopySnapshot",
        "ec2:RegisterImage",
        "ec2:Describe*"
      ],
      "Resource": "*"
    }
  ]
}
```

4. Use the **put-role-policy** command to attach the policy to the role you created. Specify the full path of the **role-policy.json** file. For example:

```
$ aws iam put-role-policy --role-name vmimport --policy-name vmimport --policy-document
file:///home/sample/ImportService/role-policy.json
```

Additional resources

- [VM Import Service Role](#)
- [Required Service Role](#)

3.4.4. Converting and pushing an image to S3

By using the **qemu-img** command, you can convert your image, so that you can push it to S3. The samples are representative; they convert an image formatted in the **qcow2** file format to **raw** format. Amazon accepts images in **OVA**, **VHD**, **VHDX**, **VMDK**, and **raw** formats. See [How VM Import/Export Works](#) for more information about image formats that Amazon accepts.

Procedure

1. Run the **qemu-img** command to convert your image. For example:

```
# qemu-img convert -f qcow2 -O raw rhel-10.0-sample.qcow2 rhel-{ProductNumber}.0-sample.raw
```

2. Push the image to S3.

```
$ aws s3 cp rhel-10.0-sample.raw s3://s3-bucket-name
```



NOTE

This procedure could take a few minutes. After completion, you can check that your image uploaded successfully to your S3 bucket by using the [AWS S3 Console](#).

Additional resources

- [How VM Import/Export Works](#)
- [AWS S3 Console](#)

3.4.5. Importing an image as a snapshot

To launch a RHEL instance in the Amazon Elastic Cloud Compute (EC2) service, you require an Amazon Machine Image (AMI). To create an AMI of your system, you must first upload a snapshot of your RHEL system image to EC2.

Procedure

1. Create a file to specify a bucket and path for your image. Name the file **containers.json**. In the sample that follows, replace **s3-bucket-name** with your bucket name and **s3-key** with your key. You can get the key for the image by using the Amazon S3 Console.

```
{
  "Description": "rhel-10.0-sample.raw",
  "Format": "raw",
  "UserBucket": {
    "S3Bucket": "s3-bucket-name",
    "S3Key": "s3-key"
  }
}
```

2. Import the image as a snapshot. This example uses a public Amazon S3 file; you can use the [Amazon S3 Console](#) to change permissions settings on your bucket.

```
$ aws ec2 import-snapshot --disk-container file://containers.json
```

The terminal displays a message such as the following. Note the **ImportTaskID** within the message.

```
{
  "SnapshotTaskDetail": {
    "Status": "active",
    "Format": "RAW",
    "DiskImageSize": 0.0,
    "UserBucket": {
      "S3Bucket": "s3-bucket-name",
      "S3Key": "rhel-10.0-sample.raw"
    },
    "Progress": "3",
    "StatusMessage": "pending"
  },
  "ImportTaskId": "import-snap-06cea01fa0f1166a8"
}
```

3. Track the progress of the import by using the **describe-import-snapshot-tasks** command. Include the **ImportTaskID**.

```
$ aws ec2 describe-import-snapshot-tasks --import-task-ids import-snap-06cea01fa0f1166a8
```

The returned message shows the current status of the task. When complete, **Status** shows **completed**. Within the status, note the snapshot ID.

Additional resources

- [Amazon S3 Console](#)
- [Importing a Disk as a Snapshot Using VM Import/Export](#)

3.4.6. Creating an AMI from the uploaded snapshot

To launch a RHEL instance in Amazon Elastic Cloud Compute (EC2) service, you require an Amazon Machine Image (AMI). To create an AMI of your system, you can use a RHEL system snapshot that you previously uploaded.

Procedure

1. Go to the AWS EC2 Dashboard.
2. Under **Elastic Block Store**, select **Snapshots**.
3. Search for your snapshot ID (for example, **snap-0e718930bd72bcdad0**).
4. Right-click on the snapshot and select **Create image**.
5. Name your image.
6. Under **Virtualization type**, choose **Hardware-assisted virtualization**.
7. Click **Create**. In the note regarding image creation, there is a link to your image.
8. Click on the image link. Your image shows up under **Images>AMIs**.

**NOTE**

Alternatively, you can use the AWS CLI **register-image** command to create an AMI from a snapshot. See [register-image](#) for more information. An example follows.

```
$ aws ec2 register-image \
  --name "myimagename" --description "myimagedescription" --architecture
x86_64 \
  --virtualization-type hvm --root-device-name "/dev/sda1" --ena-support \
  --block-device-mappings "{\"DeviceName\": \"/dev/sda1\", \"Ebs\":
  {\"SnapshotId\": \"snap-0ce7f009b69ab274d\"}}"
```

You must specify the root device volume **/dev/sda1** as your **root-device-name**. For conceptual information about device mapping for AWS, see [Example block device mapping](#).

3.4.7. Launching an instance from the AMI

To launch and configure an Amazon Elastic Compute Cloud (EC2) instance, use an Amazon Machine Image (AMI).

Procedure

1. From the AWS EC2 Dashboard, select **Images** and then **AMIs**.
2. Right-click on your image and select **Launch**.
3. Choose an **Instance Type** that meets or exceeds the requirements of your workload. See [Amazon EC2 Instance Types](#) for information about instance types.
4. Click **Next: Configure Instance Details**.
 - a. Enter the **Number of instances** you want to create.
 - b. For **Network**, select the VPC you created when [setting up your AWS environment](#). Select a subnet for the instance or create a new subnet.
 - c. Select **Enable** for Auto-assign Public IP.

**NOTE**

These are the minimum configuration options necessary to create a basic instance. Review additional options based on your application requirements.

5. Click **Next: Add Storage**. Verify that the default storage is sufficient.
6. Click **Next: Add Tags**.

**NOTE**

Tags can help you manage your AWS resources. See [Tagging Your Amazon EC2 Resources](#) for information about tagging.

7. Click **Next: Configure Security Group**. Select the security group you created when [setting up your AWS environment](#).
8. Click **Review and Launch**. Verify your selections.
9. Click **Launch**. You are prompted to select an existing key pair or create a new key pair. Select the key pair you created when [setting up your AWS environment](#).

**NOTE**

Verify that the permissions for your private key are correct. Use the command options **chmod 400 <keyname>.pem** to change the permissions, if necessary.

10. Click **Launch Instances**.
11. Click **View Instances**. You can name the instance(s).
You can now launch an SSH session to your instance(s) by selecting an instance and clicking **Connect**. Use the example provided for **A standalone SSH client**

**NOTE**

Alternatively, you can launch an instance by using the AWS CLI. See [Launching, Listing, and Terminating Amazon EC2 Instances](#) in the Amazon documentation for more information.

Additional resources

- [AWS Management Console](#)
- [Setting Up with Amazon EC2](#)
- [Amazon EC2 Instances](#)
- [Amazon EC2 Instance Types](#)

3.4.8. Attaching Red Hat subscriptions

Using the **subscription-manager** command, you can register and attach your Red Hat subscription to a RHEL instance.

Prerequisites

- You have an active [Red Hat account](#).

Procedure

1. Register your system:

```
# subscription-manager register --auto-attach
```

2. Attach your subscriptions:

- You can use an activation key to attach subscriptions. See [Creating Red Hat Customer Portal Activation Keys](#) for more information.

- Alternatively, you can manually attach a subscription by using the ID of the subscription pool (Pool ID). See [Attaching a host-based subscription to hypervisors](#).
3. Optional: To collect various system metrics about the instance in the [Red Hat Hybrid Cloud Console](#), you can register the instance with [Red Hat Insights](#).

```
# insights-client register --display-name <display-name-value>
```

For information on further configuration of Red Hat Insights, see [Client Configuration Guide for Red Hat Insights](#).

Additional resources

- [Creating Red Hat Customer Portal Activation Keys](#)
- [Attaching a host-based subscription to hypervisors](#)
- [Client Configuration Guide for Red Hat Insights](#)

3.4.9. Setting up automatic registration on AWS Gold Images

To make deploying RHEL virtual machines on Amazon Web Services (AWS) faster and more comfortable, you can set up Gold Images of RHEL to be automatically registered to the Red Hat Subscription Manager (RHSM).

Prerequisites

- You have downloaded the latest RHEL Gold Image for AWS. For instructions, see [Using Gold Images on AWS](#).



NOTE

An AWS account can only be attached to a single Red Hat account at a time. Therefore, ensure no other users require access to the AWS account before attaching it to your Red Hat one.

Procedure

1. Upload the Gold Image to AWS. For instructions, see [Uploading the Red Hat Enterprise Linux image to AWS](#).
2. Create VMs by using the uploaded image. They will be automatically subscribed to RHSM.

Verification

- In a RHEL VM created using the above instructions, verify the system is registered to RHSM by executing the **subscription-manager identity** command. On a successfully registered system, this displays the UUID of the system. For example:

```
# subscription-manager identity
system identity: fdc46662-c536-43fb-a18a-bbcb283102b7
name: 192.168.122.222
org name: 6340056
org ID: 6340056
```


Additional resources

- [AWS Management Console](#)
- [Adding cloud integrations to the Hybrid Cloud Console](#)

3.5. ADDITIONAL RESOURCES

- [Red Hat Cloud Access Reference Guide](#)
- [Red Hat in the Public Cloud](#)
- [Red Hat Enterprise Linux on Amazon EC2 - FAQs](#)
- [Setting Up with Amazon EC2](#)
- [Red Hat on Amazon Web Services](#)

CHAPTER 4. CONFIGURING A RED HAT HIGH AVAILABILITY CLUSTER ON AWS

To create a cluster where RHEL nodes automatically redistribute their workloads if a node failure occurs, use the Red Hat High Availability Add-On. Such high availability (HA) clusters can also be hosted on public cloud platforms, including AWS. Creating RHEL HA clusters on AWS is similar to creating HA clusters in non-cloud environments.

To configure a Red Hat HA cluster on Amazon Web Services (AWS) using EC2 instances as cluster nodes, see the following sections. Note that you have a number of options for obtaining the RHEL images you use for your cluster. For information on image options for AWS, see [Red Hat Enterprise Linux Image Options on AWS](#).

Prerequisites

- Sign up for a [Red Hat Customer Portal](#) account.
- Sign up for AWS and set up your AWS resources. See [Setting Up with Amazon EC2](#) for more information.

4.1. BENEFITS OF USING HIGH-AVAILABILITY CLUSTERS ON PUBLIC CLOUD PLATFORMS

A high-availability (HA) cluster is a set of computers (called *nodes*) that are linked together to run a specific workload. The purpose of HA clusters is to provide redundancy in case of a hardware or software failure. If a node in the HA cluster fails, the Pacemaker cluster resource manager distributes the workload to other nodes and no noticeable downtime occurs in the services that are running on the cluster.

You can also run HA clusters on public cloud platforms. In this case, you would use virtual machine (VM) instances in the cloud as the individual cluster nodes. Using HA clusters on a public cloud platform has the following benefits:

- *Improved availability:* In case of a VM failure, the workload is quickly redistributed to other nodes, so running services are not disrupted.
- *Scalability:* Additional nodes can be started when demand is high and stopped when demand is low.
- *Cost-effectiveness:* With the pay-as-you-go pricing, you pay only for nodes that are running.
- *Simplified management:* Some public cloud platforms offer management interfaces to make configuring HA clusters easier.

To enable HA on your RHEL systems, Red Hat offers a High Availability Add-On. The High Availability Add-On provides all necessary components for creating HA clusters on RHEL systems. The components include high availability service management and cluster administration tools.

Additional resources

- [High Availability Add-On overview](#)

4.2. CREATING THE AWS ACCESS KEY AND AWS SECRET ACCESS KEY

You need to create an AWS Access Key and AWS Secret Access Key before you install the AWS CLI. The fencing and resource agent APIs use the AWS Access Key and Secret Access Key to connect to each node in the cluster.

Prerequisites

- Your IAM user account must have Programmatic access. See [Setting up the AWS Environment](#) for more information.

Procedure

1. Launch the [AWS Console](#).
2. Click on your AWS Account ID to display the drop-down menu and select **My Security Credentials**.
3. Click **Users**.
4. Select the user and open the **Summary** screen.
5. Click the **Security credentials** tab.
6. Click **Create access key**.
7. Download the **.csv** file (or save both keys). You need to enter these keys when creating the fencing device.

4.3. CREATING AN EC2 INSTANCE FOR HA CLUSTER

Complete the following steps to create the instances that you use as your HA cluster nodes. Note that you have a number of options for obtaining the RHEL images you use for your cluster. See [Red Hat Enterprise Linux Image options on AWS](#) for information about image options for AWS.

You can create and upload a custom image that you use for your cluster nodes, or you can use a Gold Image or an on-demand image.

Prerequisites

- You have set up an AWS environment. For more information, see [Setting Up with Amazon EC2](#).

Procedure

1. From the AWS EC2 Dashboard, select **Images** and then **AMIs**.
2. Right-click on your image and select **Launch**.
3. Choose an **Instance Type** that meets or exceeds the requirements of your workload. Depending on your HA application, each instance may need to have higher capacity. See [Amazon EC2 Instance Types](#) for information about instance types.
4. Click **Next: Configure Instance Details**.
 - a. Enter the **Number of instances** you want to create for the cluster. This example procedure uses three cluster nodes.

**NOTE**

Do not launch into an Auto Scaling Group.

- b. For **Network**, select the VPC you created in [Set up the AWS environment](#). Select the subnet for the instance to create a new subnet.
- c. Select **Enable** for Auto-assign Public IP. These are the minimum selections you need to make for **Configure Instance Details**. Depending on your specific HA application, you may need to make additional selections.

**NOTE**

These are the minimum configuration options necessary to create a basic instance. Review additional options based on your HA application requirements.

5. Click **Next: Add Storage** and verify that the default storage is sufficient. You do not need to modify these settings unless your HA application requires other storage options.
6. Click **Next: Add Tags**.

**NOTE**

Tags can help you manage your AWS resources. See [Tagging Your Amazon EC2 Resources](#) for information about tagging.

7. Click **Next: Configure Security Group**. Select the existing security group you created in [Setting up the AWS environment](#).
8. Click **Review and Launch** and verify your selections.
9. Click **Launch**. You are prompted to select an existing key pair or create a new key pair. Select the key pair you created when [Setting up the AWS environment](#).
10. Click **Launch Instances**.
11. Click **View Instances**. You can name the instance(s).

**NOTE**

Alternatively, you can launch instances by using the AWS CLI. See [Launching, Listing, and Terminating Amazon EC2 Instances](#) in the Amazon documentation for more information.

Additional resources

- [AWS Management Console](#)
- [Setting Up with Amazon EC2](#)
- [Amazon EC2 Instances](#)
- [Amazon EC2 Instance Types](#)

4.4. CONFIGURING THE PRIVATE KEY

Complete the following configuration tasks to use the private SSH key file (**.pem**) before it can be used in an SSH session.

Procedure

1. Move the key file from the **Downloads** directory to your **Home** directory or to your **~/.ssh directory**.
2. Change the permissions of the key file so that only the root user can read it.

```
# chmod 400 KeyName.pem
```

4.5. CONNECTING TO AN EC2 INSTANCE

Using the **AWS Console** on all nodes, you can connect to an EC2 instance.

Procedure

1. Launch the [AWS Console](#) and select the EC2 instance.
2. Click **Connect** and select **A standalone SSH client**.
3. From your SSH terminal session, connect to the instance by using the AWS example provided in the pop-up window. Add the correct path to your **KeyName.pem** file if the path is not shown in the example.

4.6. INSTALLING THE HIGH AVAILABILITY PACKAGES AND AGENTS

On each of the nodes, you need to install the High Availability packages and agents to be able to configure a Red Hat High Availability cluster on AWS.

Procedure

1. Remove the AWS Red Hat Update Infrastructure (RHUI) client.

```
$ sudo -i  
# dnf -y remove rh-amazon-rhui-client*
```

2. Register the VM with Red Hat.

```
# subscription-manager register --auto-attach
```

3. Disable all repositories.

```
# subscription-manager repos --disable=*
```

4. Enable the RHEL 10 Server HA repositories.

```
# subscription-manager repos --enable=rhel-10-for-x86_64-highavailability-rpms
```

5. Update the RHEL AWS instance.

```
# dnf update -y
```

6. Install the Red Hat High Availability Add-On software packages, along with the AWS fencing agent from the High Availability channel.

```
# dnf install pcs pacemaker fence-agents-aws
```

7. The user **hacluster** was created during the **pcs** and **pacemaker** installation in the previous step. Create a password for **hacluster** on all cluster nodes. Use the same password for all nodes.

```
# passwd hacluster
```

8. Add the **high availability** service to the RHEL Firewall if **firewalld.service** is installed.

```
# firewall-cmd --permanent --add-service=high-availability
# firewall-cmd --reload
```

9. Start the **pcs** service and enable it to start on boot.

```
# systemctl start pcsd.service
# systemctl enable pcsd.service
```

10. Edit **/etc/hosts** and add RHEL host names and internal IP addresses. See [How should the /etc/hosts file be set up on RHEL cluster nodes?](#) for details.

Verification

- Ensure the **pcs** service is running.

```
# systemctl status pcsd.service
```

```
pcsd.service - PCS GUI and remote configuration interface
Loaded: loaded (/usr/lib/systemd/system/pcsd.service; enabled; vendor preset: disabled)
Active: active (running) since Thu 2018-03-01 14:53:28 UTC; 28min ago
Docs: man:pcsd(8)
man:pcs(8)
Main PID: 5437 (pcsd)
CGroup: /system.slice/pcsd.service
└─5437 /usr/bin/ruby /usr/lib/pcsd/pcsd > /dev/null &
Mar 01 14:53:27 ip-10-0-0-48.ec2.internal systemd[1]: Starting PCS GUI and remote
configuration interface...
Mar 01 14:53:28 ip-10-0-0-48.ec2.internal systemd[1]: Started PCS GUI and remote
configuration interface.
```

4.7. CREATING A CLUSTER

Complete the following steps to create the cluster of nodes.

Procedure

1. On one of the nodes, enter the following command to authenticate the pcs user **hacluster**. In the command, specify the name of each node in the cluster.

```
# pcs host auth <hostname1> <hostname2> <hostname3>
```

Example:

```
[root@node01 clouduser]# pcs host auth node01 node02 node03
Username: hacluster
Password:
node01: Authorized
node02: Authorized
node03: Authorized
```

2. Create the cluster.

```
# pcs cluster setup <cluster_name> <hostname1> <hostname2> <hostname3>
```

Example:

```
[root@node01 clouduser]# pcs cluster setup new_cluster node01 node02 node03

[...]

Synchronizing pcsd certificates on nodes node01, node02, node03...
node02: Success
node03: Success
node01: Success
Restarting pcsd on the nodes in order to reload the certificates...
node02: Success
node03: Success
node01: Success
```

Verification

1. Enable the cluster.

```
[root@node01 clouduser]# pcs cluster enable --all
node02: Cluster Enabled
node03: Cluster Enabled
node01: Cluster Enabled
```

2. Start the cluster.

```
[root@node01 clouduser]# pcs cluster start --all
node02: Starting Cluster...
node03: Starting Cluster...
node01: Starting Cluster...
```

4.8. CONFIGURING FENCING

Fencing configuration ensures that a malfunctioning node on your AWS cluster is automatically isolated, which prevents the node from consuming the cluster's resources or compromising the cluster's functionality.

To configure fencing on an AWS cluster, you can use multiple methods:

- A standard procedure for default configuration.
- An alternate configuration procedure for more advanced configuration, focused on automation.

Prerequisites

- You must be using the **fence_aws** fencing agent. To obtain **fence_aws**, install the **resource-agents** package on your cluster.

Standard procedure

1. Enter the following AWS metadata query to get the Instance ID for each node. You need these IDs to configure the fence device. See [Instance Metadata and User Data](#) for additional information.

```
# echo $(curl -s http://169.254.169.254/latest/meta-data/instance-id)
```

Example:

```
[root@ip-10-0-0-48 ~]# echo $(curl -s http://169.254.169.254/latest/meta-data/instance-id) i-07f1ac63af0ec0ac6
```

2. Enter the following command to configure the fence device. Use the **pcmk_host_map** command to map the RHEL host name to the Instance ID. Use the AWS Access Key and AWS Secret Access Key that you previously set up.

```
# pcs stonith \
  create <name> fence_aws access_key=access-key secret_key=<secret-access-key> \
  region=<region> pcmk_host_map="rhel-hostname-1:Instance-ID-1;rhel-hostname-2:Instance-ID-2;rhel-hostname-3:Instance-ID-3" \
  power_timeout=240 pcmk_reboot_timeout=480 pcmk_reboot_retries=4
```

Example:

```
[root@ip-10-0-0-48 ~]# pcs stonith \
  create clusterfence fence_aws access_key=AKIAI123456MRMJA
  secret_key=a75EYIG4RVL3hdsdAslK7koQ8dzaDyn5yoIZ/ \
  region=us-east-1 pcmk_host_map="ip-10-0-0-48:i-07f1ac63af0ec0ac6;ip-10-0-0-46:i-063fc5fe93b4167b2;ip-10-0-0-58:i-08bd39eb03a6fd2c7" \
  power_timeout=240 pcmk_reboot_timeout=480 pcmk_reboot_retries=4
```

Alternate procedure

1. Obtain the VPC ID of the cluster.

```
# aws ec2 describe-vpcs --output text --filters "Name=tag:Name,Values=<clustername>-vpc"
--query 'Vpcs[0].VpcId'
vpc-06bc10ac8f6006664
```


- By using the VPC ID of the cluster, obtain the VPC instances.

```
$ aws ec2 describe-instances --output text --filters "Name=vpc-id,Values=vpc-
06bc10ac8f6006664" --query 'Reservations[*].Instances[*].{Name:Tags[? Key==Name][
0].Value,Instance:InstanceId}' | grep "\-node[a-c]"
i-0b02af8927a895137    <clustername>-nodea-vm
i-0cceb4ba8ab743b69    <clustername>-nodeb-vm
i-0502291ab38c762a5    <clustername>-nodec-vm
```

- Use the obtained instance IDs to configure fencing on each node on the cluster. For example, to configure a fencing device on all nodes in a cluster:

```
[root@nodea ~]# CLUSTER=<clustername> && pcs stonith create fence${CLUSTER}
fence_aws access_key=XXXXXXXXXXXXXXXXXXXXX pcmk_host_map=$(for NODE \
in node{a..c}; do ssh ${NODE} "echo -n \${HOSTNAME}:\$(curl -s
http://169.254.169.254/latest/meta-data/instance-id)\;"; done) \
pcmk_reboot_retries=4 pcmk_reboot_timeout=480 power_timeout=240 region=xx-xxxx-x
secret_key=XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
```

For information about specific parameters for creating fencing devices, see the **fence_aws** man page or the [Configuring and managing high availability clusters](#) guide.

Verification

- Display the configured fencing devices and their parameters on your nodes:

```
[root@nodea ~]# pcs stonith config fence${CLUSTER}

Resource: <clustername> (class=stonith type=fence_aws)
Attributes: access_key=XXXXXXXXXXXXXXXXXXXXX pcmk_host_map=nodea:i-
0b02af8927a895137;nodeb:i-0cceb4ba8ab743b69;nodec:i-0502291ab38c762a5;
pcmk_reboot_retries=4 pcmk_reboot_timeout=480 power_timeout=240 region=xx-xxxx-x
secret_key=XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
Operations: monitor interval=60s (<clustername>-monitor-interval-60s)
```

- Test the fencing agent for one of the cluster nodes.

```
# pcs stonith fence <awsnodename>
```



NOTE

The command response may take several minutes to display. If you watch the active terminal session for the node being fenced, you see that the terminal connection is immediately terminated after you enter the fence command.

Example:

```
[root@ip-10-0-0-48 ~]# pcs stonith fence ip-10-0-0-58

Node: ip-10-0-0-58 fenced
```

- Check the status to verify that the node is fenced.

```
# pcs status
```

Example:

```
[root@ip-10-0-0-48 ~]# pcs status

Cluster name: newcluster
Stack: corosync
Current DC: ip-10-0-0-46 (version 1.1.18-11.el7-2b07d5c5a9) - partition with quorum
Last updated: Fri Mar  2 19:55:41 2018
Last change: Fri Mar  2 19:24:59 2018 by root via cibadmin on ip-10-0-0-46

3 nodes configured
1 resource configured

Online: [ ip-10-0-0-46 ip-10-0-0-48 ]
OFFLINE: [ ip-10-0-0-58 ]

Full list of resources:
clusterfence (stonith:fence_aws):  Started ip-10-0-0-46

Daemon Status:
corosync: active/disabled
pacemaker: active/disabled
pcsd: active/enabled
```

4. Start the node that was fenced in the previous step.

```
# pcs cluster start <awshostname>
```

5. Check the status to verify the node started.

```
# pcs status
```

Example:

```
[root@ip-10-0-0-48 ~]# pcs status

Cluster name: newcluster
Stack: corosync
Current DC: ip-10-0-0-46 (version 1.1.18-11.el7-2b07d5c5a9) - partition with quorum
Last updated: Fri Mar  2 20:01:31 2018
Last change: Fri Mar  2 19:24:59 2018 by root via cibadmin on ip-10-0-0-48

3 nodes configured
1 resource configured

Online: [ ip-10-0-0-46 ip-10-0-0-48 ip-10-0-0-58 ]

Full list of resources:

clusterfence (stonith:fence_aws):  Started ip-10-0-0-46

Daemon Status:
```

```
corosync: active/disabled
pacemaker: active/disabled
pcsd: active/enabled
```

4.9. VERIFYING CLUSTER NODES CONFIGURATION

Previously, you installed the AWS CLI on your host system. You need to install the AWS CLI on cluster nodes before you configure the network resource agents.

Complete the following procedure on each cluster node.

Prerequisites

- You must have created an AWS Access Key and AWS Secret Access Key. See [Creating the AWS Access Key and AWS Secret Access Key](#) for more information.
- You have installed the [AWS CLI](#).

Procedure

- Check that the instance IDs and names are correct:
Example:

```
[root@ip-10-0-0-48 ~]# aws ec2 describe-instances --output text --query
'Reservations[*].Instances[*].[InstanceId,Tags[?Key==Name].Value]'
i-07f1ac63af0ec0ac6
ip-10-0-0-48
i-063fc5fe93b4167b2
ip-10-0-0-46
i-08bd39eb03a6fd2c7
ip-10-0-0-58
```

4.10. SETTING UP IP ADDRESS RESOURCES ON AWS

To ensure that clients that use IP addresses to access resources managed by the cluster over the network can access the resources if a failover occurs, the cluster must include *IP address resources*, which use specific network resource agents.

The RHEL HA Add-On provides a set of resource agents, which create IP address resources to manage various types of IP addresses on AWS. To decide which resource agent to configure, consider the type of AWS IP addresses in the HA cluster management. It includes following ways to create a cluster resource for managing an IP address:

- Exposed to the internet: Use the **awseip** network resource.
- Limited to a single AWS Availability Zone (AZ): Use the **awsvip** and **IPaddr2** network resources.
- Reassigns to multiple AWS AZs within the same AWS region: Use the **aws-vpc-move-ip** network resource.

**NOTE**

If the HA cluster does not manage any IP addresses, the resource agents for managing virtual IP addresses on AWS are not required. If you need further guidance for your specific deployment, consult with your AWS provider.

4.10.1. Creating an IP address resource to manage an IP address exposed to the internet

To ensure that high-availability (HA) clients can access a RHEL 10 node that uses public-facing internet connections, configure an *AWS Secondary Elastic IP Address* (**awseip**) resource to use an elastic IP address.

Prerequisites

- You have a previously configured cluster.
- Your cluster nodes must have access to the RHEL HA repositories. For more information, see [Installing the High Availability packages and agents](#).
- You have set up the AWS CLI. For instructions, see [Installing the AWS CLI](#).

Procedure

1. Install the **resource-agents-cloud** package.

```
# dnf install resource-agents-cloud
```

2. Using the AWS command-line interface (CLI), create an elastic IP address.

```
[root@ip-10-0-0-48 ~]# aws ec2 allocate-address --domain vpc --output text
eipalloc-4c4a2c45 vpc 35.169.153.122
```

3. Optional: Display the description of **awseip**. This shows the options and default operations for this agent.

```
# pcs resource describe awseip
```

4. Create the Secondary Elastic IP address resource that uses the allocated IP address that you previously specified using the AWS CLI. In addition, create a resource group that the Secondary Elastic IP address will belong to.

```
# pcs resource create <resource-id> awseip elastic_ip=<Elastic-IP-Address>
allocation_id=<Elastic-IP-Association-ID> --group networking-group
```

Example:

```
# pcs resource create elastic awseip elastic_ip=35.169.153.122 allocation_id=eipalloc-
4c4a2c45 --group networking-group
```

Verification

1. Display the status of the cluster to verify that the required resources are running.

```
# pcs status
```

The following output shows an example running cluster where the **vip** and **elastic** resources have been started as a part of the **networking-group** resource group:

```
[root@ip-10-0-0-58 ~]# pcs status

Cluster name: newcluster
Stack: corosync
Current DC: ip-10-0-0-58 (version 1.1.18-11.el7-2b07d5c5a9) - partition with quorum
Last updated: Mon Mar  5 16:27:55 2018
Last change: Mon Mar  5 15:57:51 2018 by root via cibadmin on ip-10-0-0-46

3 nodes configured
4 resources configured

Online: [ ip-10-0-0-46 ip-10-0-0-48 ip-10-0-0-58 ]

Full list of resources:

clusterfence (stonith:fence_aws): Started ip-10-0-0-46
Resource Group: networking-group
  vip (ocf::heartbeat:IPaddr2): Started ip-10-0-0-48
  elastic (ocf::heartbeat:awseip): Started ip-10-0-0-48

Daemon Status:
  corosync: active/disabled
  pacemaker: active/disabled
  pcsd: active/enabled
```

2. Launch an SSH session from your local workstation to the elastic IP address that you previously created.

```
$ ssh -l <user-name> -i ~/.ssh/<KeyName>.pem <elastic-IP>
```

Example:

```
$ ssh -l ec2-user -i ~/.ssh/cluster-admin.pem 35.169.153.122
```

3. Verify that the host to which you connected via SSH is the host associated with the elastic resource created.

4.10.2. Creating an IP address resource to manage a private IP address limited to a single AWS Availability Zone

To ensure that high-availability (HA) clients on AWS can access a RHEL node that uses a private IP address that can only move in a single AWS Availability Zone (AZ), configure an *AWS Secondary Private IP Address* (**awsvip**) resource to use a virtual IP address.

You can complete the following procedure on any node in the cluster.

Prerequisites

- You have a previously configured cluster.
- Your cluster nodes have access to the RHEL HA repositories. For more information, see [Installing the High Availability packages and agents](#).

Procedure

1. Install the **resource-agents-cloud** package.

```
# dnf install resource-agents-cloud
```

2. Optional: View the **awsvip** description. This shows the options and default operations for this agent.

```
# pcs resource describe awsvip
```

3. Create a Secondary Private IP address with an unused private IP address in the **VPC CIDR** block. In addition, create a resource group that the Secondary Private IP address will belong to.

```
# pcs resource create <resource-id> awsvip secondary_private_ip=<Unused-IP-Address> --group <group-name>
```

Example:

```
[root@ip-10-0-0-48 ~]# pcs resource create privip awsvip secondary_private_ip=10.0.0.68 --group networking-group
```

4. Create a virtual IP resource. This is a VPC IP address that can be rapidly remapped from the fenced node to the failover node, masking the failure of the fenced node within the subnet. Ensure that the virtual IP belongs to the same resource group as the Secondary Private IP address you created in the previous step.

```
# pcs resource create <resource-id> IPaddr2 ip=<secondary-private-IP> --group <group-name>
```

Example:

```
root@ip-10-0-0-48 ~]# pcs resource create vip IPaddr2 ip=10.0.0.68 --group networking-group
```

Verification

- Display the status of the cluster to verify that the required resources are running.

```
# pcs status
```

The following output shows an example running cluster where the **vip** and **privip** resources have been started as a part of the **networking-group** resource group:

```
[root@ip-10-0-0-48 ~]# pcs status
Cluster name: newcluster
Stack: corosync
Current DC: ip-10-0-0-46 (version 1.1.18-11.el7-2b07d5c5a9) - partition with quorum
```

```

Last updated: Fri Mar 2 22:34:24 2018
Last change: Fri Mar 2 22:14:58 2018 by root via cibadmin on ip-10-0-0-46

3 nodes configured
3 resources configured

Online: [ ip-10-0-0-46 ip-10-0-0-48 ip-10-0-0-58 ]

Full list of resources:

clusterfence (stonith:fence_aws): Started ip-10-0-0-46
Resource Group: networking-group
  privip (ocf::heartbeat:awsvip): Started ip-10-0-0-48
  vip (ocf::heartbeat:IPaddr2): Started ip-10-0-0-58

Daemon Status:
corosync: active/disabled
pacemaker: active/disabled
pcsd: active/enabled

```

4.10.3. Creating an IP address resource to manage an IP address that can move across multiple AWS Availability Zones

To ensure that high-availability (HA) clients on AWS can access a RHEL node that can be moved across multiple AWS Availability Zones within the same AWS region, configure an **aws-vpc-move-ip** resource to use an elastic IP address.

Prerequisites

- You have a previously configured cluster.
- Your cluster nodes have access to the RHEL HA repositories. For more information, see [Installing the High Availability packages and agents](#).
- An Identity and Access Management (IAM) user is configured on your cluster and has the following permissions:
 - Modify routing tables
 - Create security groups
 - Create IAM policies and roles

Procedure

1. Install the **resource-agents-cloud** package.

```
# dnf install resource-agents-cloud
```

2. Optional: View the **aws-vpc-move-ip** description. This shows the options and default operations for this agent.

```
# pcs resource describe aws-vpc-move-ip
```

3. Set up an **OverlayIPAgent** IAM policy for the IAM user.

- a. In the AWS console, navigate to **Services → IAM → Policies → Create OverlayIPAgent Policy**
- b. Input the following configuration, and change the `<region>`, `<account-id>`, and `<ClusterRouteTableID>` values to correspond with your cluster.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "Stmt1424870324000",
      "Effect": "Allow",
      "Action": "ec2:DescribeRouteTables",
      "Resource": "*"
    },
    {
      "Sid": "Stmt1424860166260",
      "Action": [
        "ec2:CreateRoute",
        "ec2:ReplaceRoute"
      ],
      "Effect": "Allow",
      "Resource": "arn:aws:ec2:<region>:<account-id>:route-
table/<ClusterRouteTableID>"
    }
  ]
}
```

4. In the AWS console, disable the **Source/Destination Check** function on all nodes in the cluster. To do this, right-click each node → **Networking → Change Source/Destination Checks**. In the pop-up message that appears, click **Yes, Disable**.
5. Create a route table for the cluster. To do so, use the following command on one node in the cluster:

```
# aws ec2 create-route --route-table-id <ClusterRouteTableID> --destination-cidr-block
<NewCIDRblockIP/NetMask> --instance-id <ClusterNodeID>
```

In the command, replace values as follows:

- **ClusterRouteTableID**: The route table ID for the existing cluster VPC route table.
 - **NewCIDRblockIP/NetMask**: A new IP address and netmask outside of the VPC classless inter-domain routing (CIDR) block. For example, if the VPC CIDR block is **172.31.0.0/16**, the new IP address/netmask can be **192.168.0.15/32**.
 - **ClusterNodeID**: The instance ID for another node in the cluster.
6. On one of the nodes in the cluster, create a **aws-vpc-move-ip** resource that uses a free IP address that is accessible to the client. The following example creates a resource named **vpcip** that uses IP **192.168.0.15**.

```
# pcs resource create vpcip aws-vpc-move-ip ip=192.168.0.15 interface=eth0
routing_table=<ClusterRouteTableID>
```


7. On all nodes in the cluster, edit the `/etc/hosts/` file, and add a line with the IP address of the newly created resource. For example:

```
192.168.0.15 vpcip
```

Verification

1. Test the failover ability of the new **aws-vpc-move-ip** resource:

```
# pcs resource move vpcip
```

2. If the failover succeeded, remove the automatically created constraint after the move of the **vpcip** resource:

```
# pcs resource clear vpcip
```

Additional resources

- [IAM users](#)

4.10.4. Additional resources

- [High Availability Add-On overview](#)
- [Configuring and managing high availability clusters](#)
- [An example of an configuration for **aws-vpc-move-ip**](#) (Red Hat Knowledgebase)

4.11. CONFIGURING SHARED BLOCK STORAGE

To create extra storage resources, you can configure shared block storage for a Red Hat High Availability cluster by using Amazon Elastic Block Storage (EBS) Multi-Attach volumes. Note that this procedure is optional, and the steps below assume three instances (a three-node cluster) with a 1 TB shared disk.

Prerequisites

- You have access to an [AWS Nitro System-based Amazon EC2 instance](#) .

Procedure

1. Create a shared block volume by using the AWS command [create-volume](#).

```
$ aws ec2 create-volume --availability-zone <availability_zone> --no-encrypted --size 1024 --volume-type io1 --iops 51200 --multi-attach-enabled
```

For example, the following command creates a volume in the **us-east-1a** availability zone.

```
$ aws ec2 create-volume --availability-zone us-east-1a --no-encrypted --size 1024 --volume-type io1 --iops 51200 --multi-attach-enabled

{
```

```

    "AvailabilityZone": "us-east-1a",
    "CreateTime": "2020-08-27T19:16:42.000Z",
    "Encrypted": false,
    "Size": 1024,
    "SnapshotId": "",
    "State": "creating",
    "VolumeId": "vol-042a5652867304f09",
    "Iops": 51200,
    "Tags": [ ],
    "VolumeType": "io1"
  }

```



NOTE

You need the **VolumeId** in the next step.

- For each instance in your cluster, attach a shared block volume by using the AWS command [attach-volume](#). Use your **<instance_id>** and **<volume_id>**.

```
$ aws ec2 attach-volume --device /dev/xvdd --instance-id <instance_id> --volume-id
<volume_id>
```

For example, the following command attaches a shared block volume **vol-042a5652867304f09** to **instance i-0eb803361c2c887f2**.

```

$ aws ec2 attach-volume --device /dev/xvdd --instance-id i-0eb803361c2c887f2 --volume-id
vol-042a5652867304f09

{
  "AttachTime": "2020-08-27T19:26:16.086Z",
  "Device": "/dev/xvdd",
  "InstanceId": "i-0eb803361c2c887f2",
  "State": "attaching",
  "VolumeId": "vol-042a5652867304f09"
}

```

Verification

- For each instance in your cluster, verify that the block device is available by using the **ssh** command with your instance **<ip_address>**.

```
# ssh <ip_address> "hostname ; lsblk -d | grep ' 1T '"
```

For example, the following command lists details including the host name and block device for the instance IP **198.51.100.3**.

```

# ssh 198.51.100.3 "hostname ; lsblk -d | grep ' 1T '"

nodea
nvme2n1 259:1  0  1T  0 disk

```

- Use the **ssh** command to verify that each instance in your cluster uses the same shared disk.

```
# ssh <ip_address> "hostname ; lsblk -d | grep ' 1T ' | awk '{print \$1}' | xargs -i udevadm info
--query=all --name=/dev/{} | grep '^E: ID_SERIAL='"
```

For example, the following command lists details including the host name and shared disk volume ID for the instance IP address **198.51.100.3**.

```
# ssh 198.51.100.3 "hostname ; lsblk -d | grep ' 1T ' | awk '{print \$1}' | xargs -i udevadm info -
-query=all --name=/dev/{} | grep '^E: ID_SERIAL='"
nodea
E: ID_SERIAL=Amazon Elastic Block Store_vol0fa5342e7aedef09f7
```

Additional resources

- [Configuring a GFS2 file system in a cluster](#)
- [Configuring GFS2 file systems](#)

4.12. ADDITIONAL RESOURCES

- [Red Hat Cloud Access Reference Guide](#)
- [Red Hat in the Public Cloud](#)
- [Red Hat on Amazon Web Services](#)
- [Red Hat Enterprise Linux on Amazon EC2 - FAQs](#)
- [Setting Up with Amazon EC2](#)